

Paderborn  
Center for  
Parallel  
Computing

# Introduction to Julia for High- Performance Computing

---

Carsten Bauer

February 13, 2024  
University College London, Advanced Research Computing Center (UCL ARC)



# Course Objectives

- ▶ Learn about **Julia** for HPC

- ▶ Learn about **HPC** with Julia

- ▶ Have some fun!

	Tuesday	Wednesday	Thursday	Friday
	Foundations	Single Core	Node	Cluster
09:00 - 10:45 (1.75h)	Julia for HPC Fundamentals	Optimising Performance I	Multithreading	Distributed Computing
10:45 - 11:00 (15m)	Break	Break	Break	Break
11:00 - 12:30 (1.5h)	Compilation	Exercises	Exercises	Exercises
12:30 - 13:30 (1h)	Lunch	Lunch	Lunch	Lunch
13:30 - 15:00 (1.5h)	Fast & Generic Code	Optimising Performance II	GPU Programming	Exercises
15:00 - 15:30 (30m)	Break	Break	Break	Q&A
15:30 - 17:00 (1.5h)	Exercises	Exercises	Exercises	

# Quick Survey

<https://etc.ch/Q7QB>



# The Julia Programming Language

## Domain science

- Dynamic & **interactive**
- High-level syntax
- Strong math support
- Package manager

## High Performance Computing

- **Compiled**
- Multithreading
- Distributed computing (e.g. MPI)
- Hardware accelerators

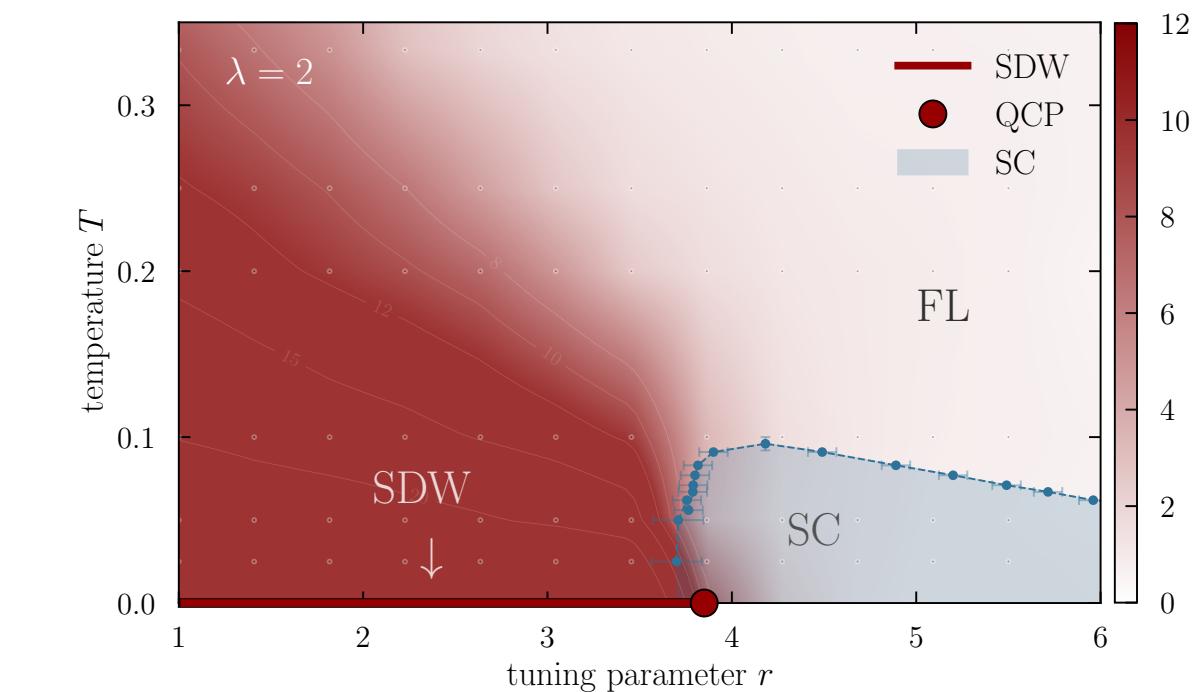
```
● ● ●
julia> function sumup()
           x = 0
           for i in 1:100
               x += i
           end
           return x
       end
sumup (generic function with 1 method)

julia> @code_llvm debuginfo=:none sumup()
define i64 @julia_sumup_177() #0 {
top:
    ret i64 5050
}
```

# Where I Come From

## Condensed matter physics / quantum field theory

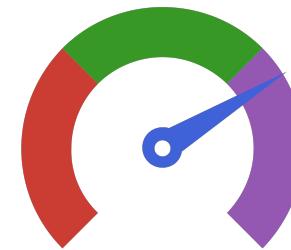
- Large-scale Quantum Monte Carlo simulations
- $O(15 \text{ Mio. CPU-core hours})$
- $O(100 \text{ TB})$



Domain Science



julia



HPC

# Exclusive HPC

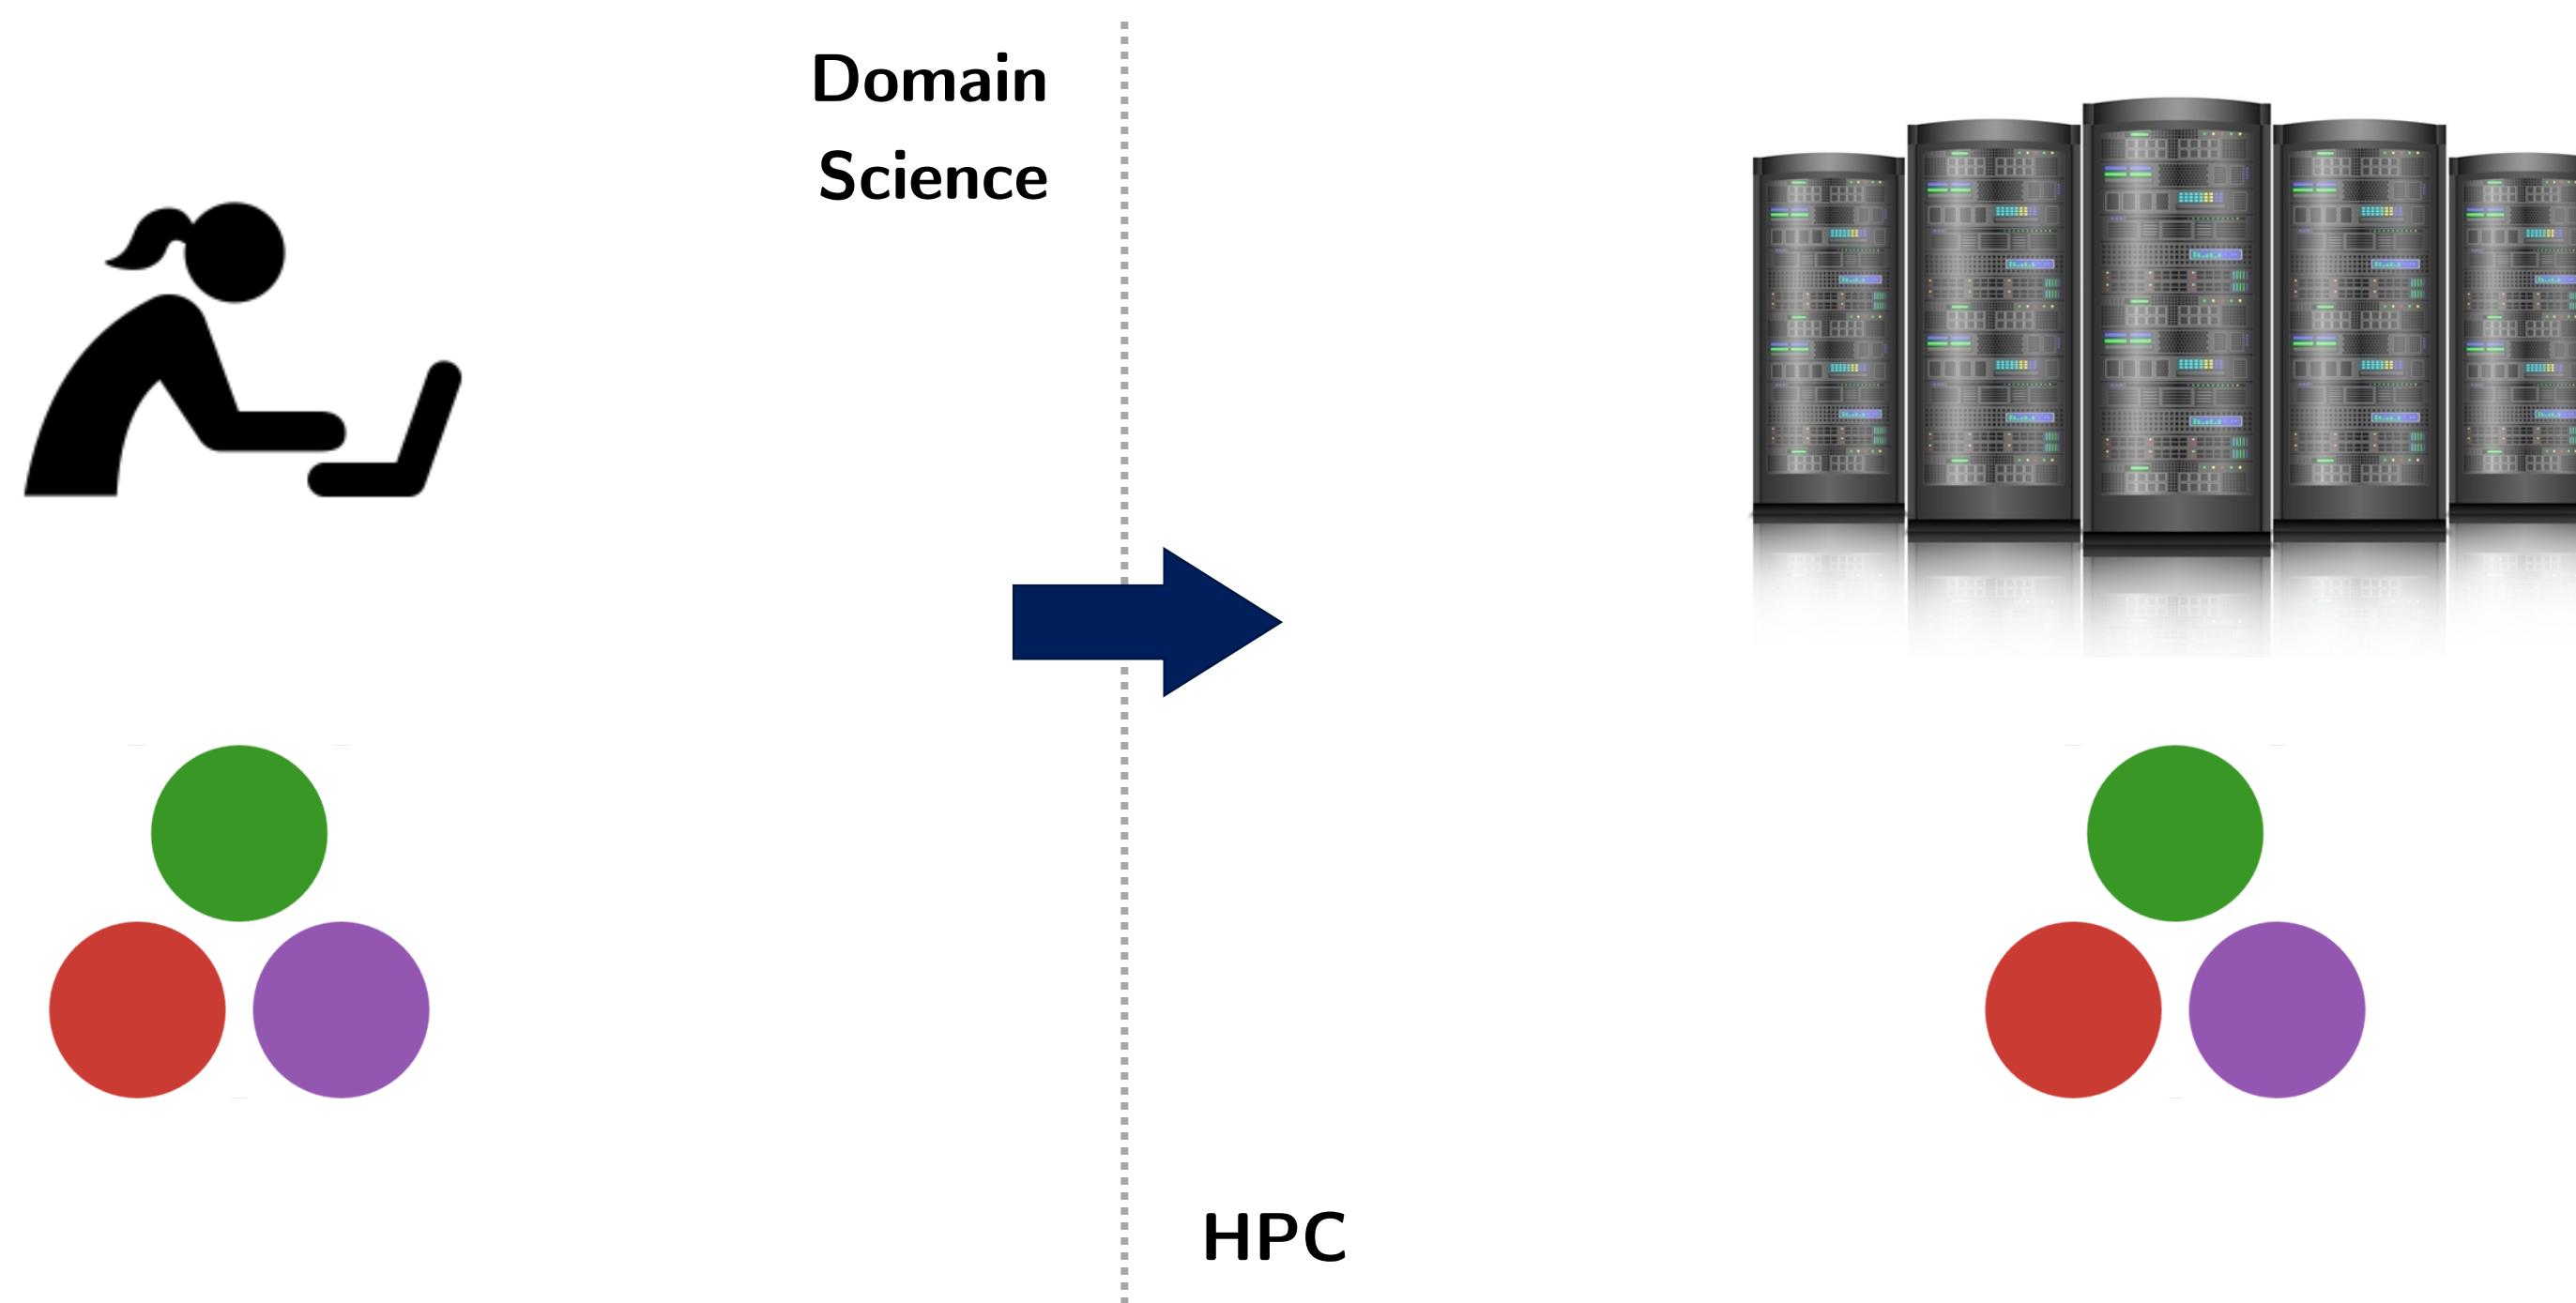
Domain  
Science



HPC



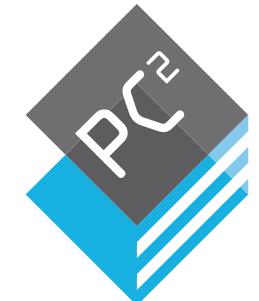
# Inclusive HPC



## Counting Flops

```
• x = rand(10_000);  
  
• function computation(x)  
•     x .+ x  
• end;
```

Counted Flops: 10000



Paderborn  
Center  
for  
Parallel  
Computing



## How?

```
• using LIKWID  
  
• metrics, events = @perfmon "FLOPS_DP" computation(x);
```

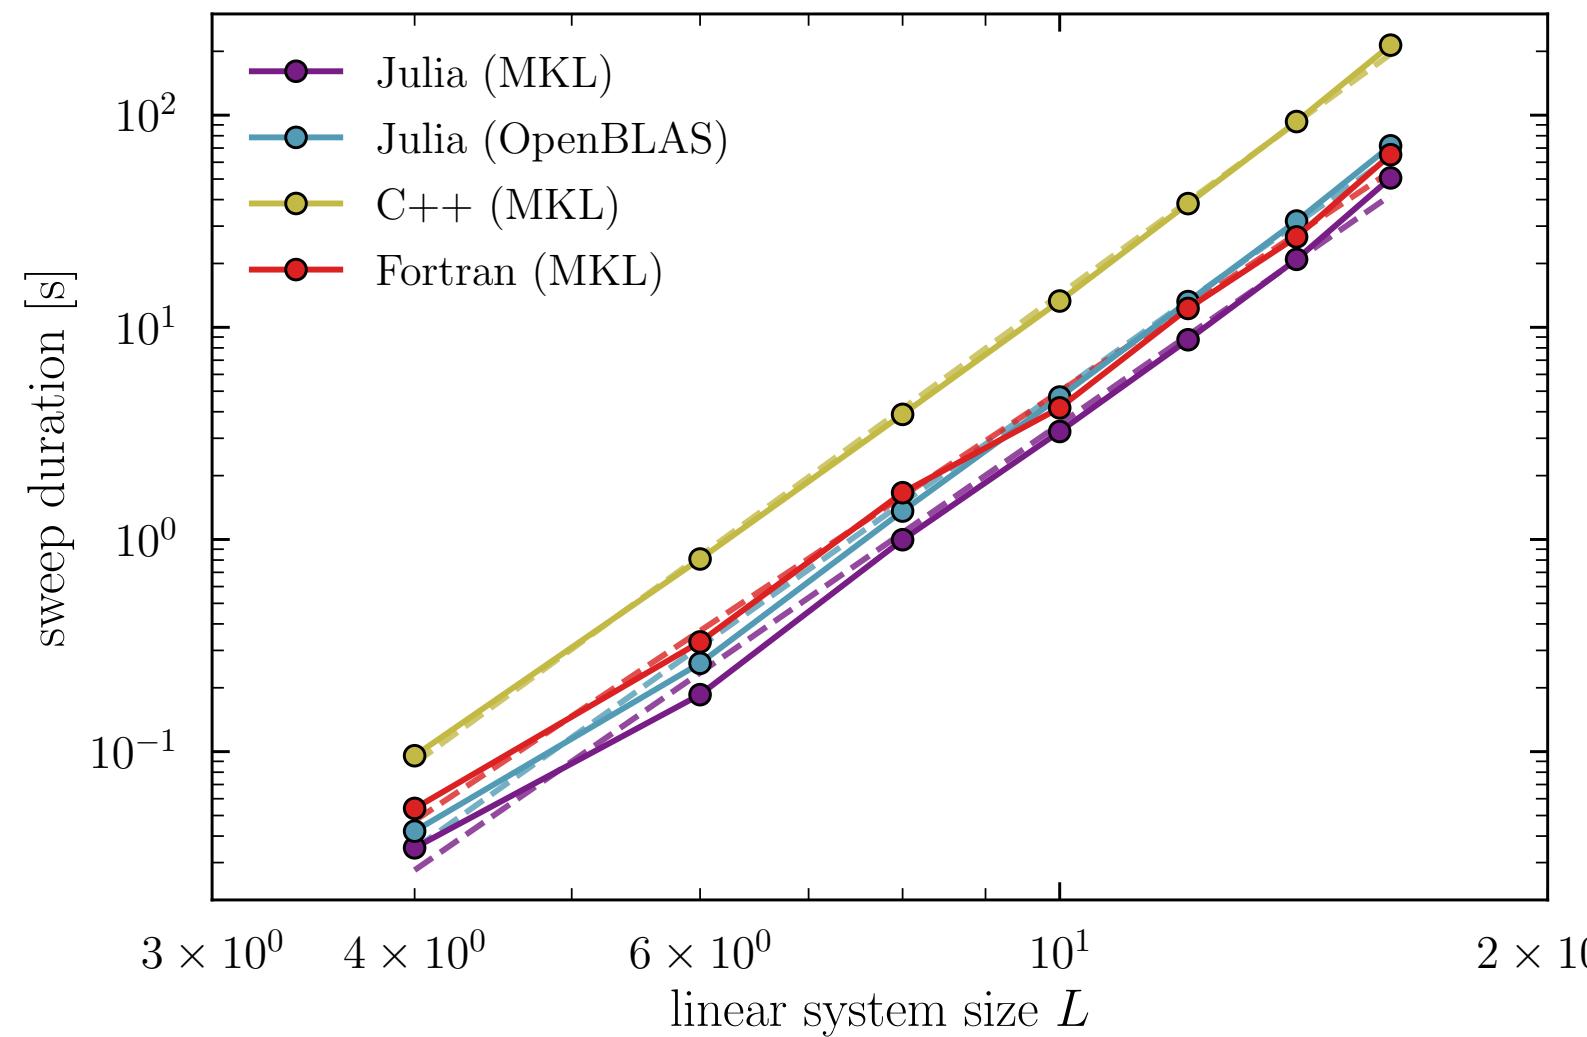
Compute from derived metrics

```
10000  
• begin  
•     flops_per_second = metrics["FLOPS_DP"][1]["DP [MFLOP/s]"] * 1e6  
•     runtime = metrics["FLOPS_DP"][1]["Runtime (RDTSC) [s]"]  
•     flops = round(Int, flops_per_second * runtime)  
• end
```

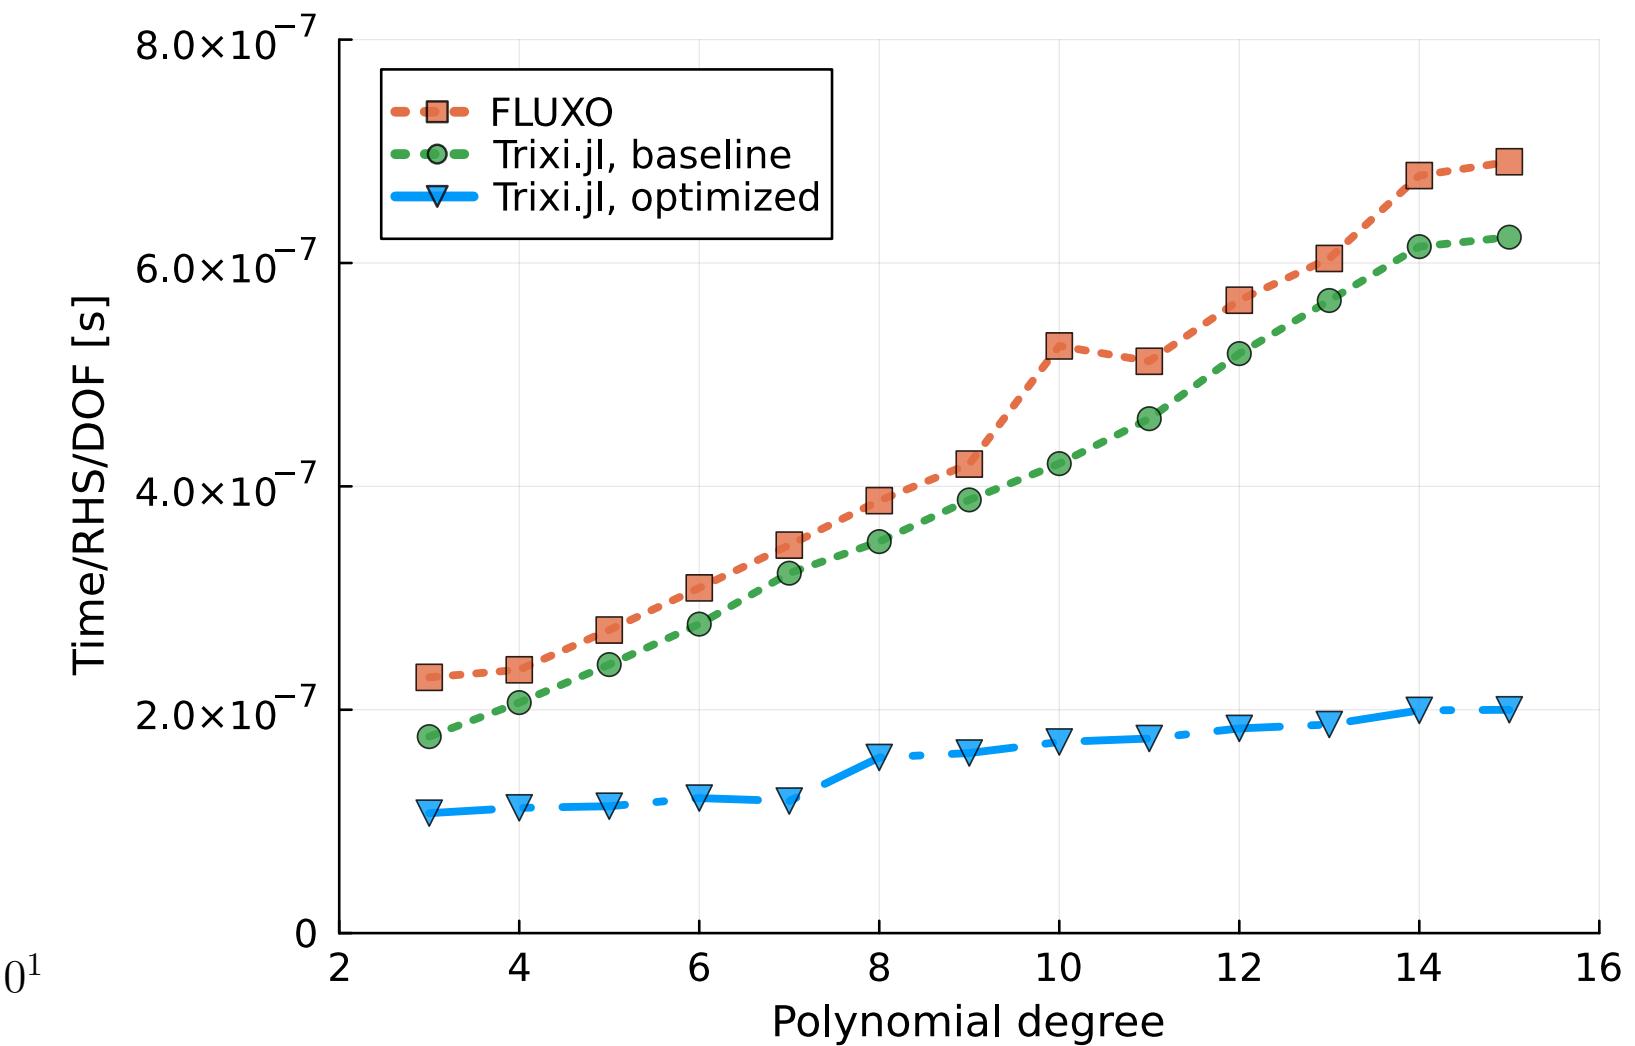


More: <https://youtu.be/I2fTNfEDPC0>

# Julia Can Be Fast



**QMC**  
(MonteCarlo.jl)



**CFD**  
(Trixi.jl)

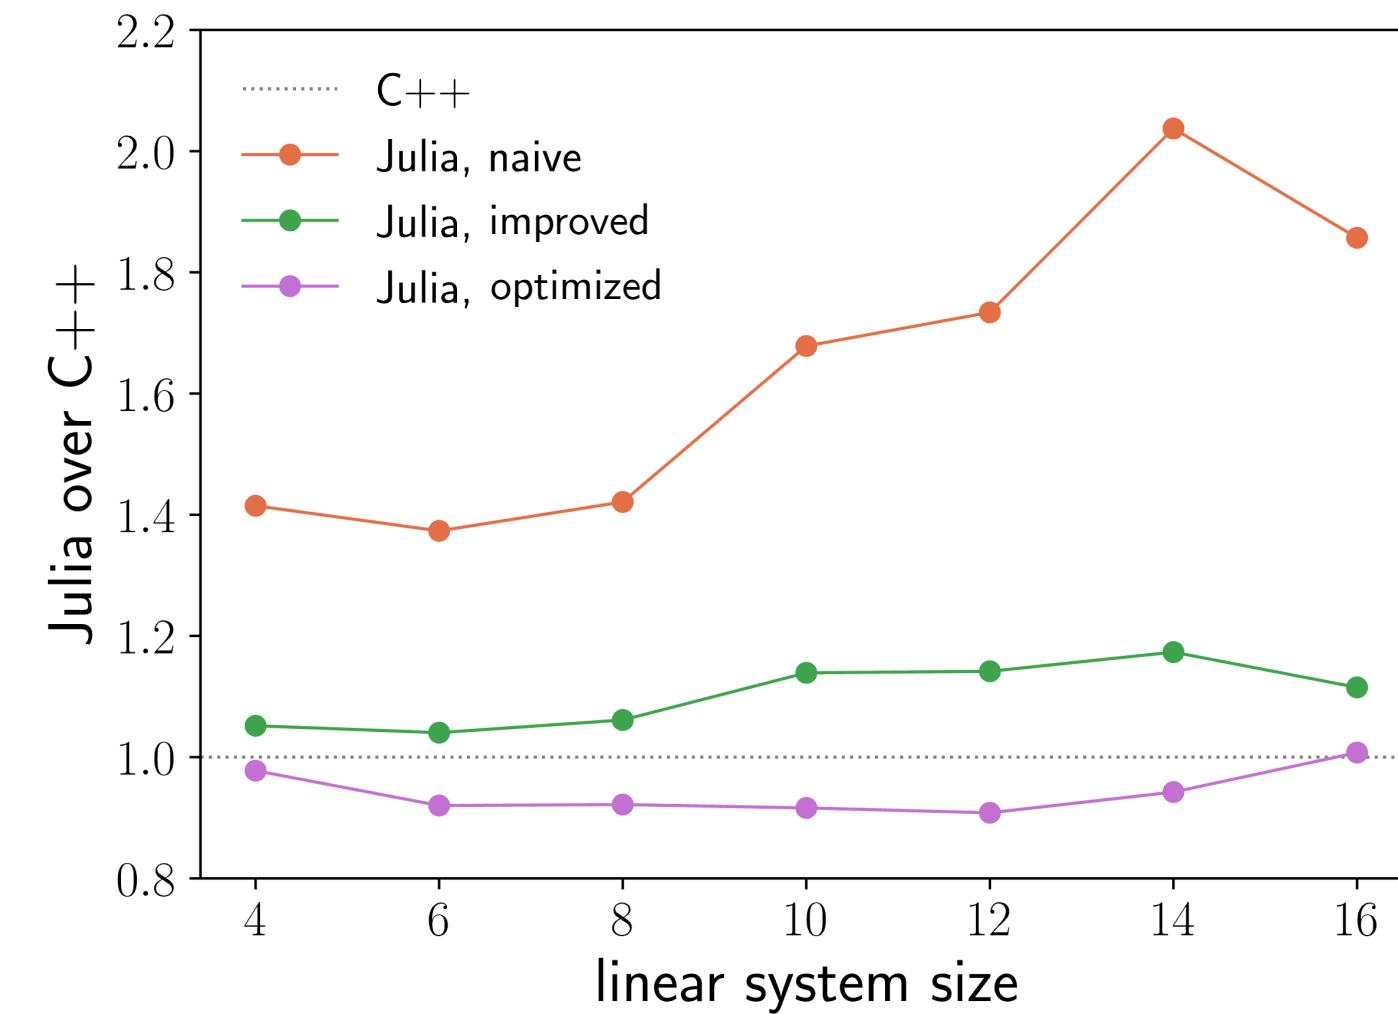
# Julia can be fast

## First Julia implementation

- ▶ ~5x faster than Python
- ▶ ~2x slower than C++

**Gradual** performance improvement, no disruptive language change!

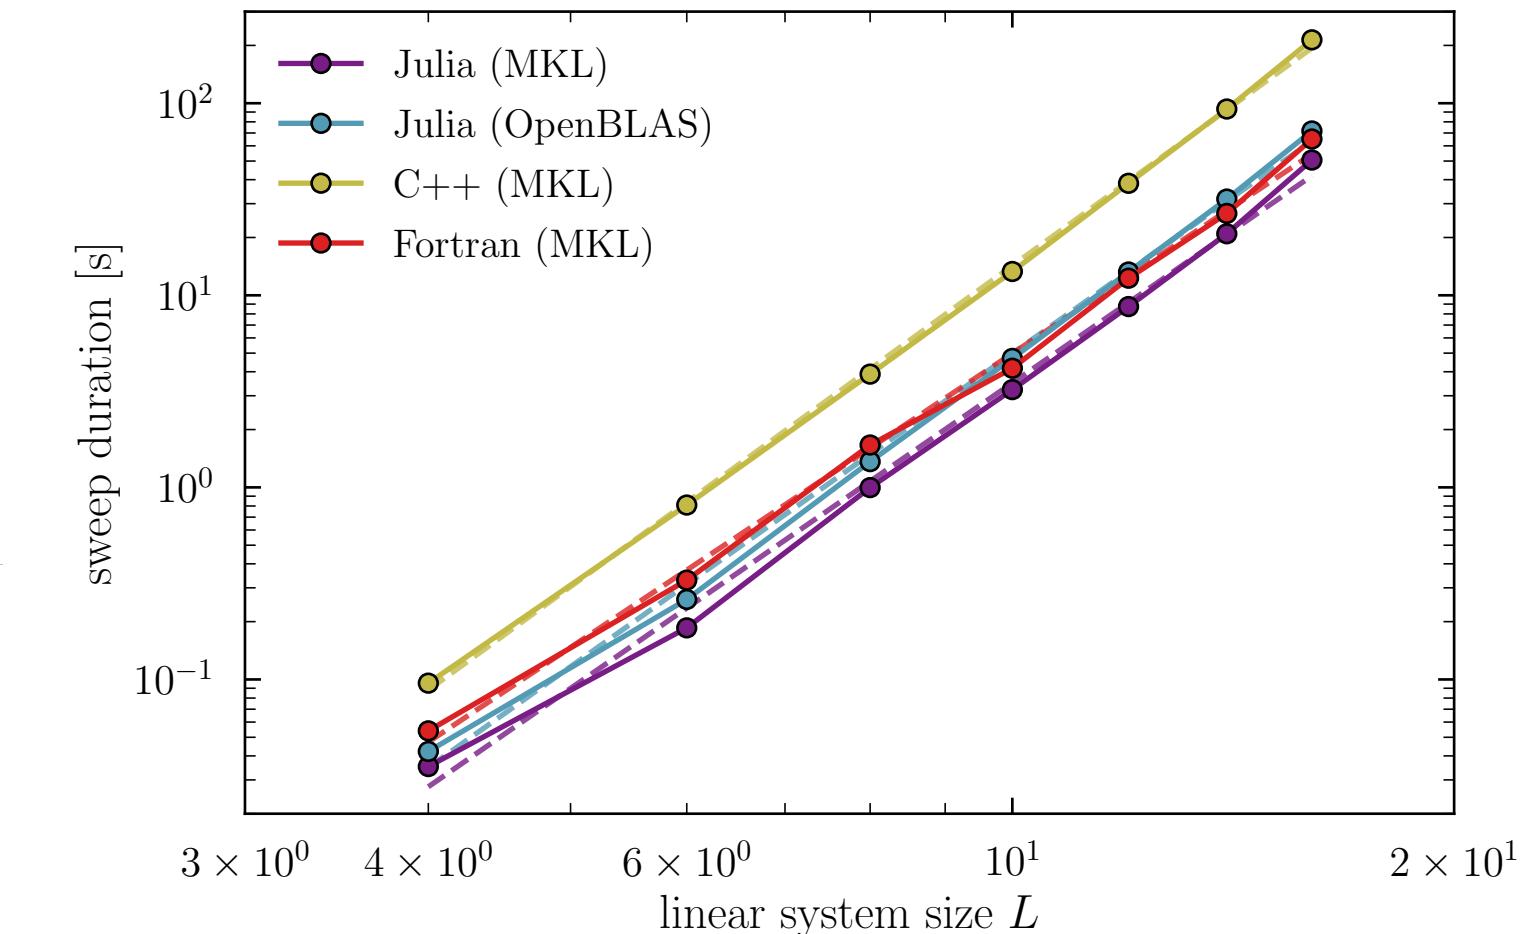
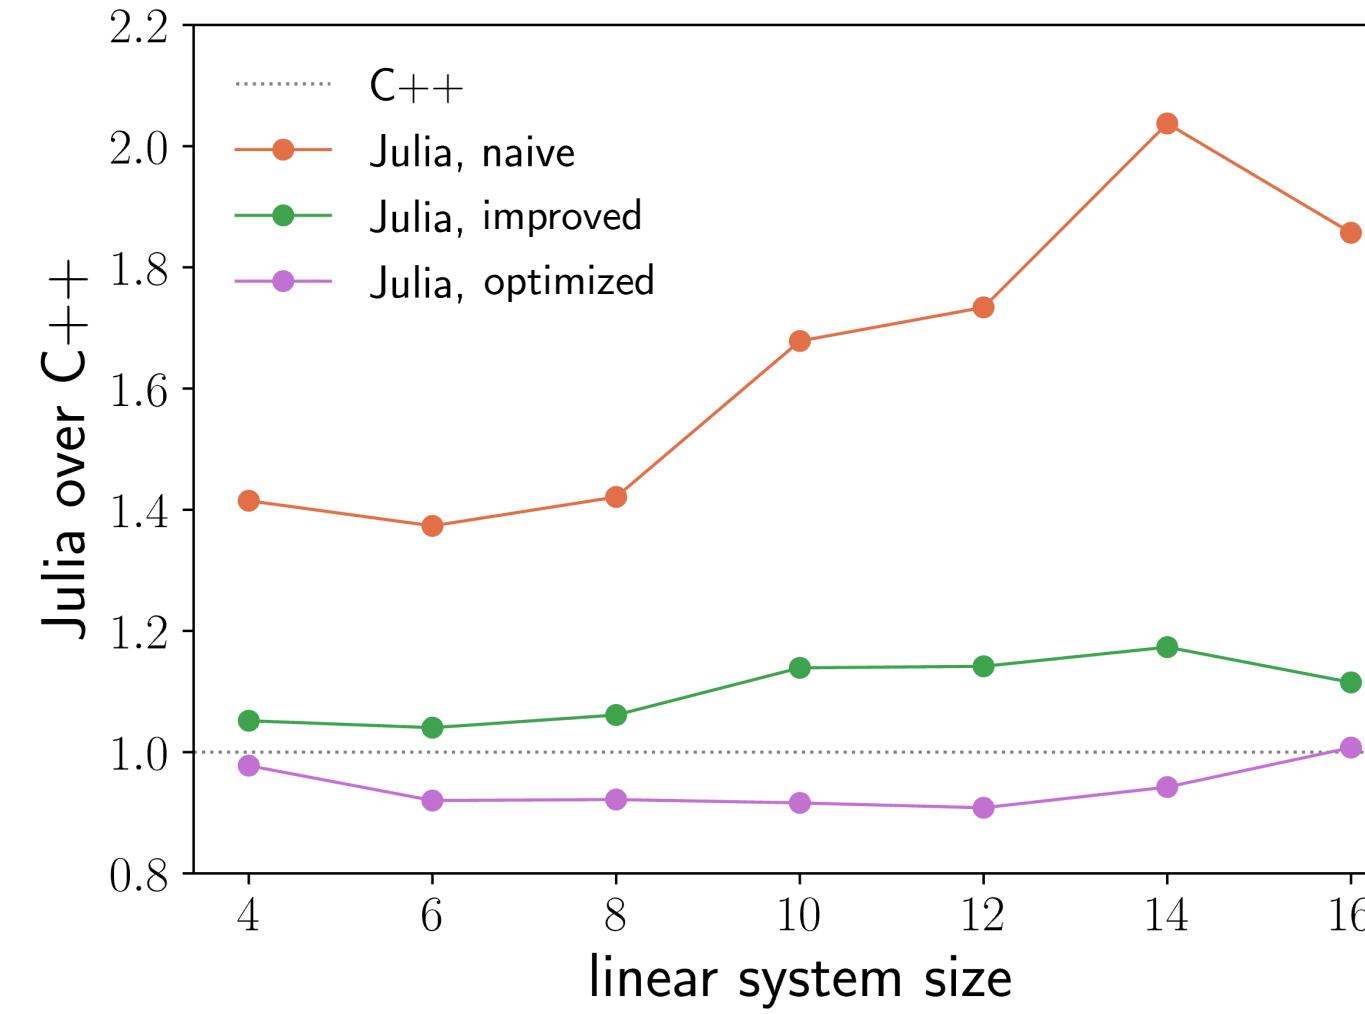
(A lot of fun!)



# Performance

**Gradual** performance improvement, no disruptive language change!

(A lot of fun!)



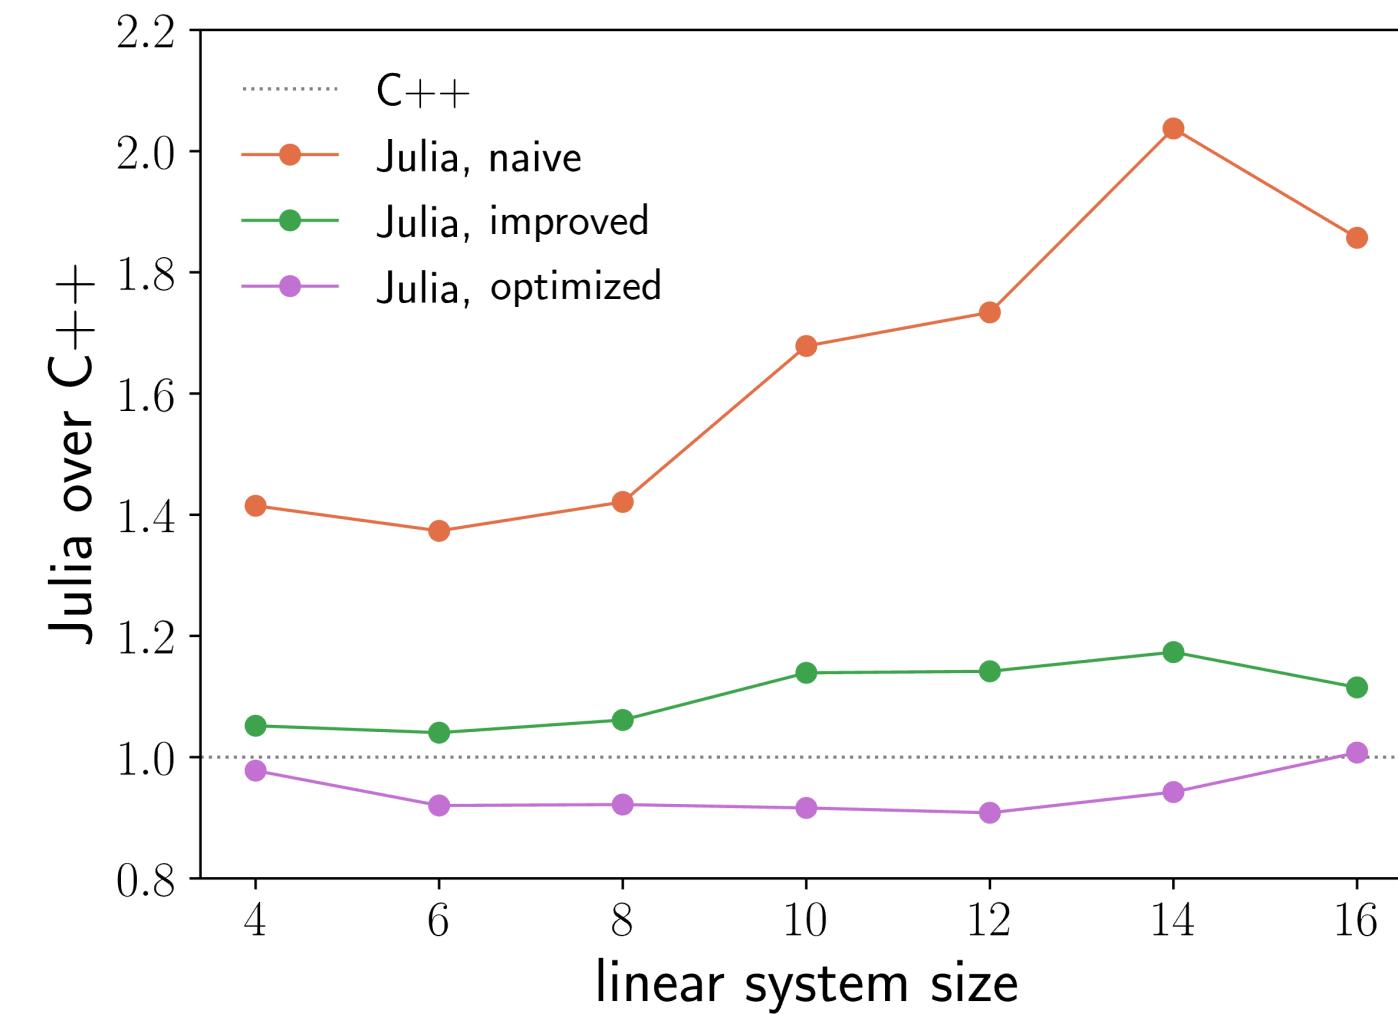
# Julia can be fast

## First Julia implementation

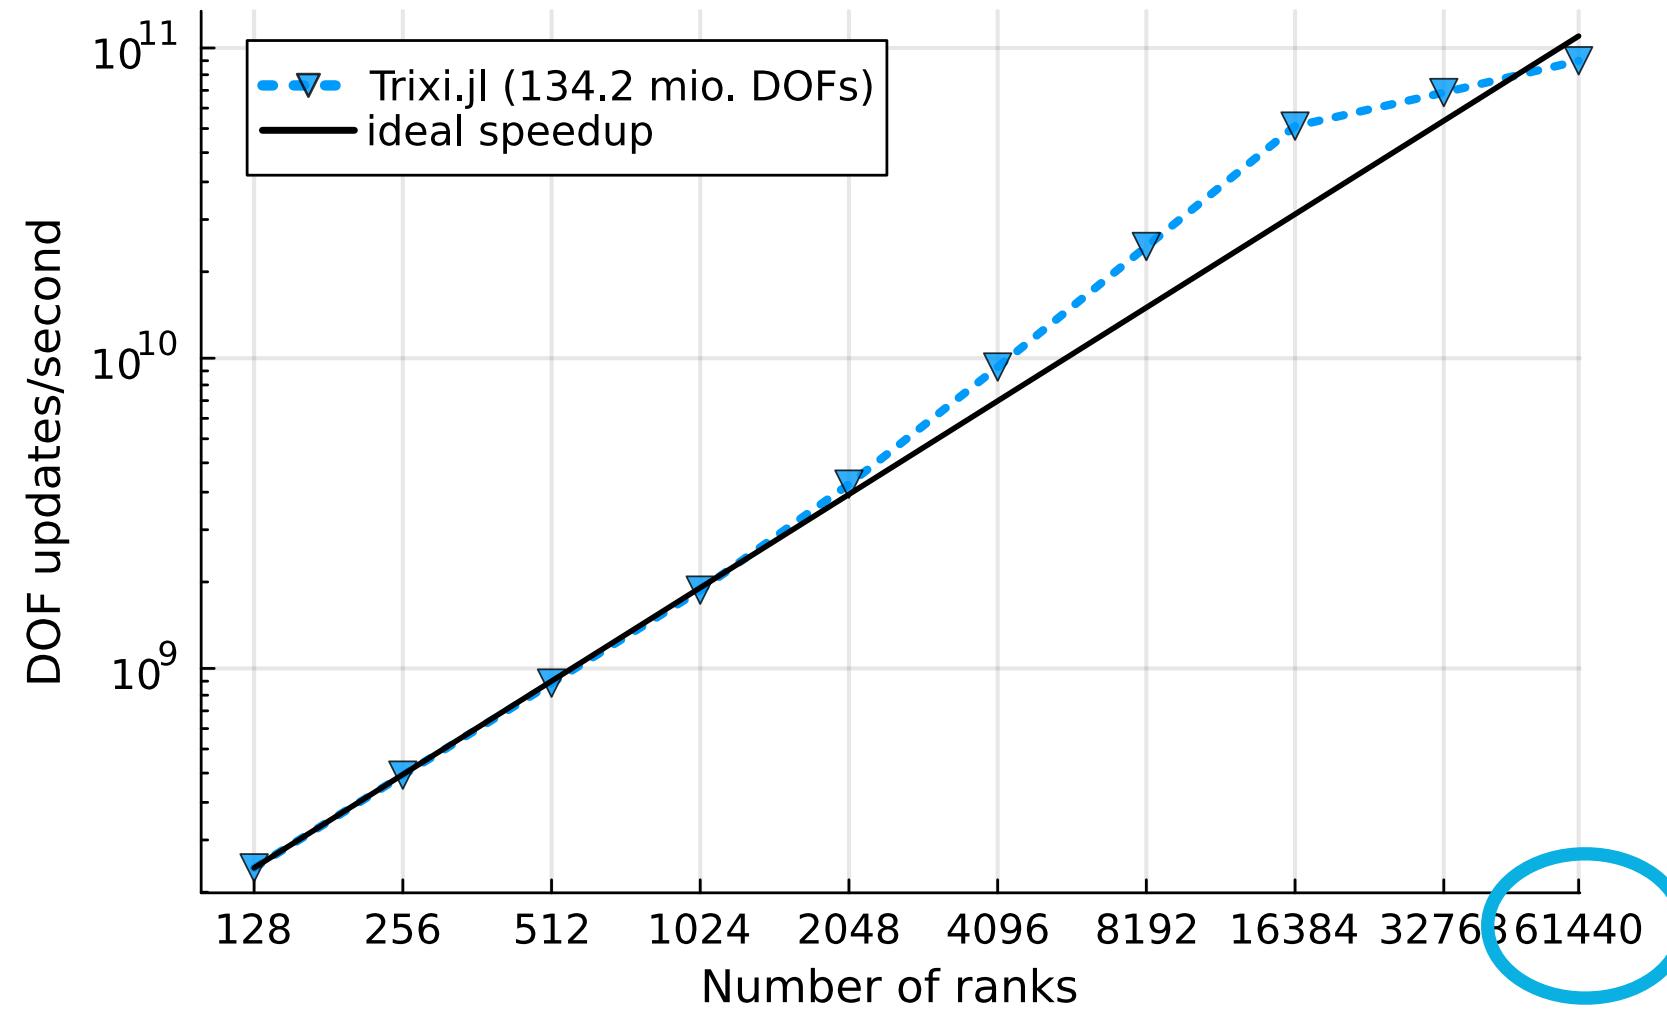
- ▶ ~5x faster than Python
- ▶ ~2x slower than C++

**Gradual** performance improvement, no disruptive language change!

(A lot of fun!)

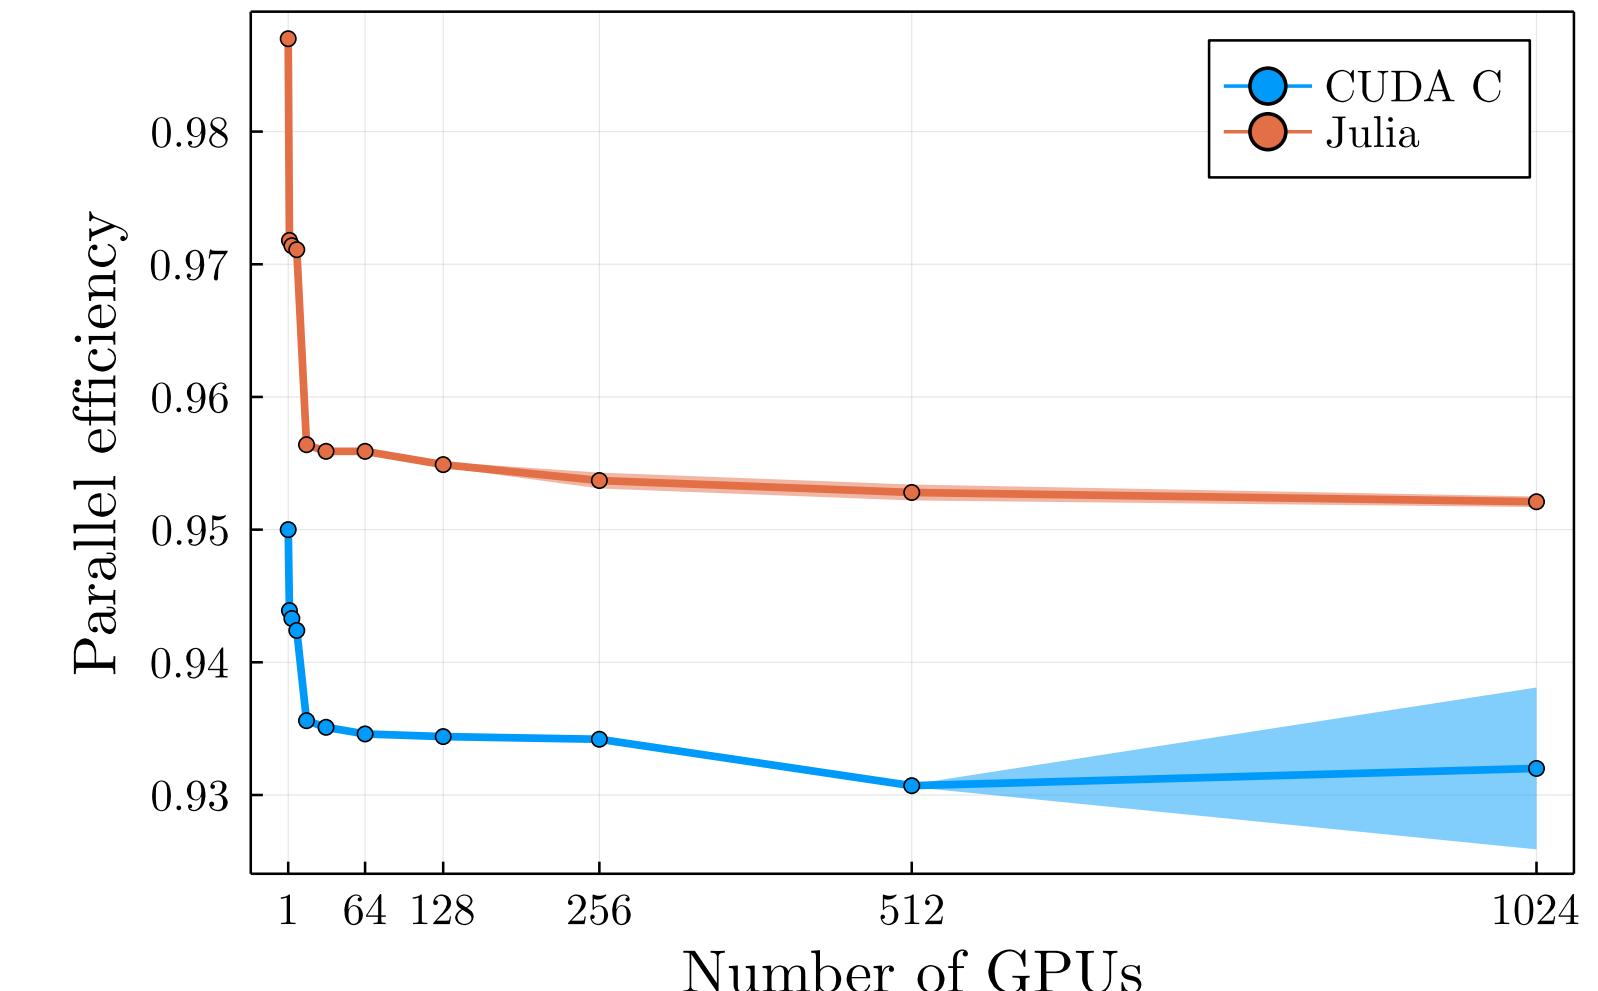


# Julia Can Be Parallel



**MPI-Parallel**  
(`Trixi.jl`)

$\approx 60\text{k ranks}$



**GPU-Parallel**  
(`ParallelStencil.jl`)

# Julia Can Be Portable

## Laptop

```
③ ⓘ ⓘ ⓘ
→ ~/myproject tree
.
└── Manifest.toml
└── Project.toml
└── code.jl

0 directories, 3 files

→ ~/myproject cat Project.toml
[deps]
CUDA = "052768ef-5323-5732-b1bb-66c8b64840ba"
DifferentialEquations = "0c46a032-eb83-5123-abaf-570d42b7fbaa"
MKL = "33e6dc65-8f57-5167-99aa-e5a354878fb2"
MPI = "da04e1cc-30fd-572f-bb4f-1f8673147195"

→ ~/myproject █
```

# HPC Cluster

(As long as you stay within the Julia-verse. Using **system software** can be a bit more tricky.)

# New Challenges (at Scale)

Dynamism comes at a price

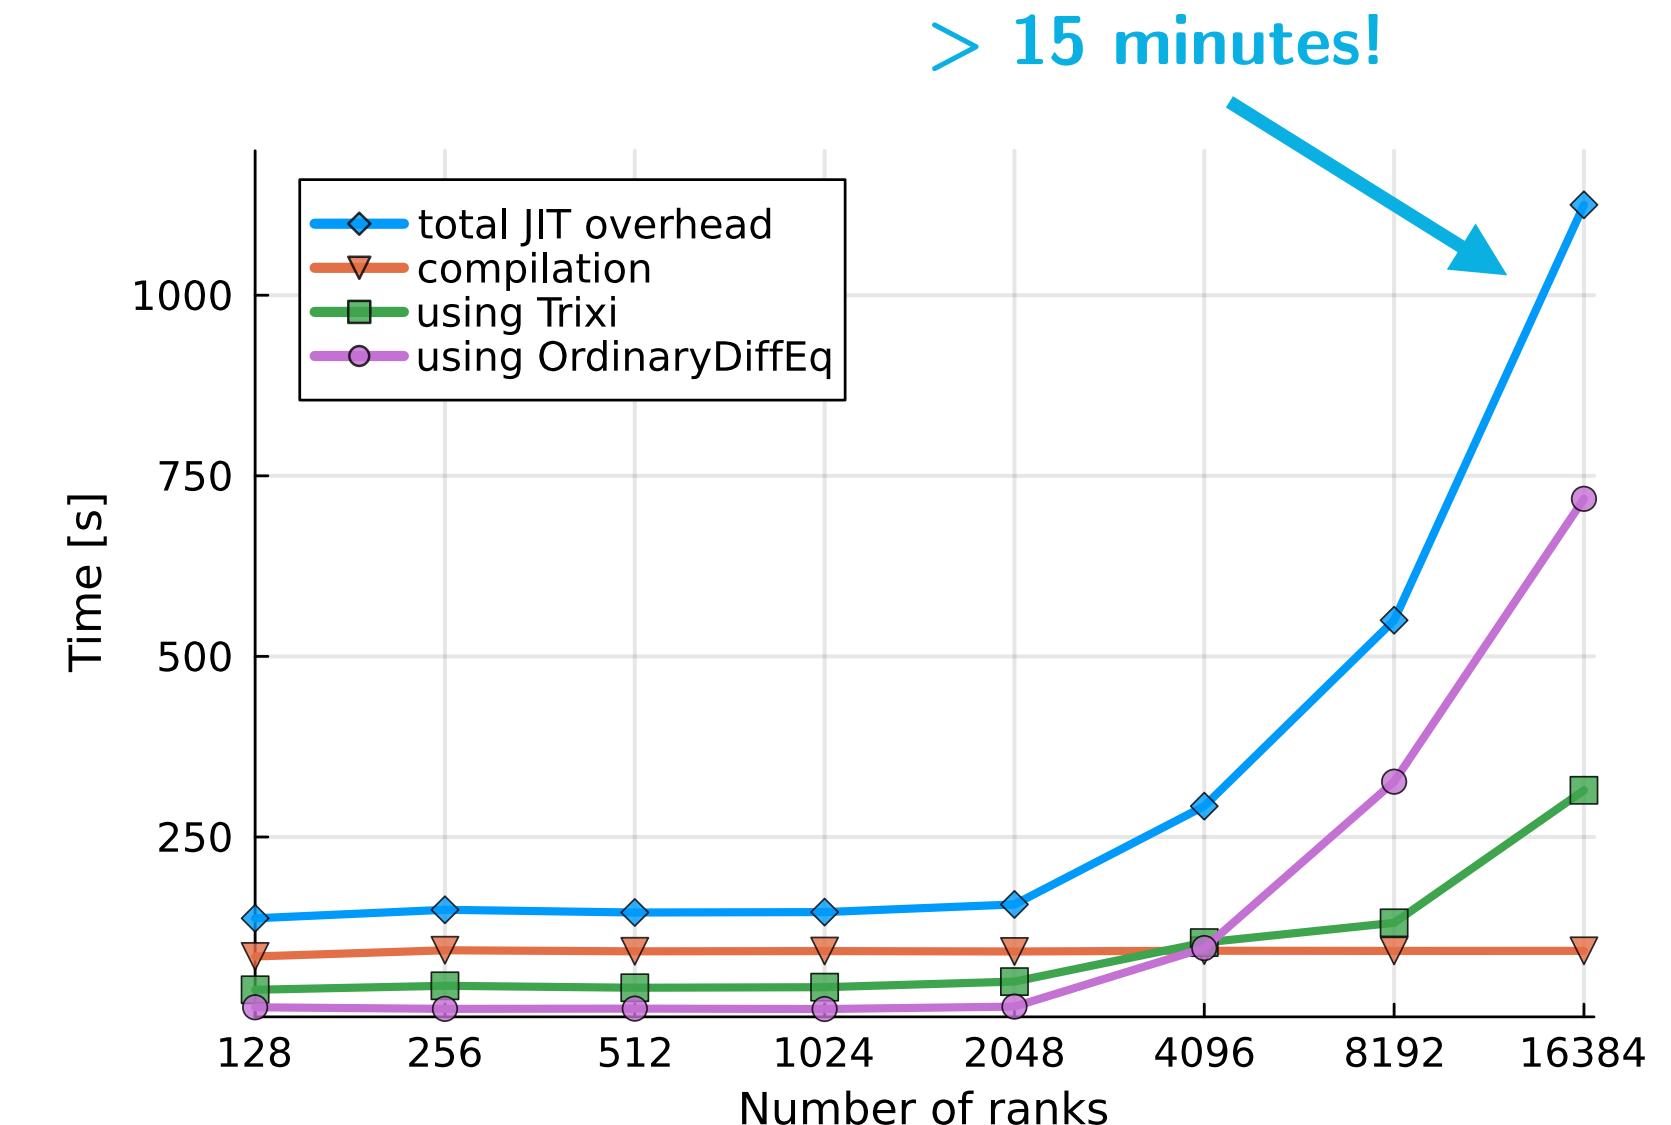
- **Type instabilities**
- No (small) **binaries**

**Julia depot** can get under (a lot) of pressure!

- Location matters (NFS vs GPFS vs RAM)

Many more challenges:

- “Interactivity at scale”
- Tooling
- ...



# Selected Julia HPC Projects

## ClIMA @ Caltech

- Climate Modeling Alliance

## CESMIX @ MIT

- Exascale simulation of materials in extreme environments

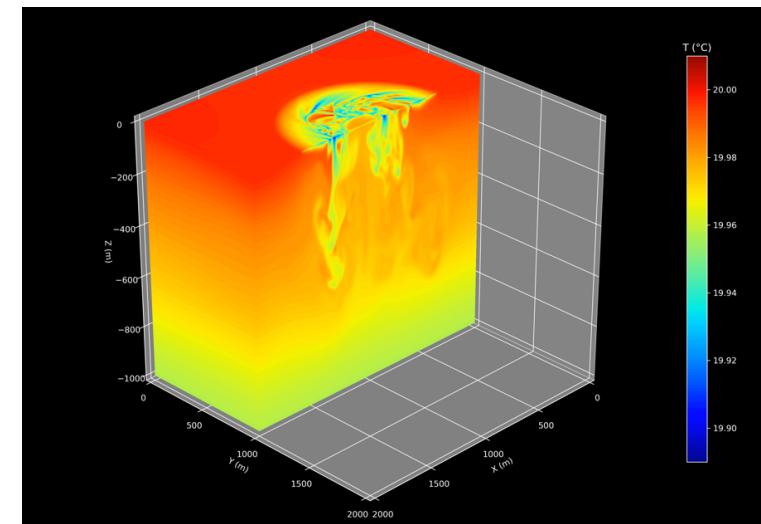
## Trixi @ RWTH Aachen / HLRS

- Computational fluid dynamics

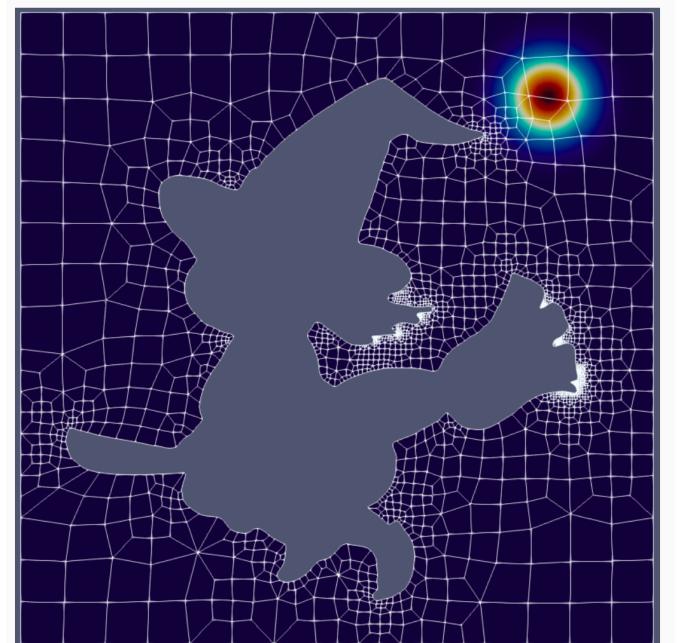
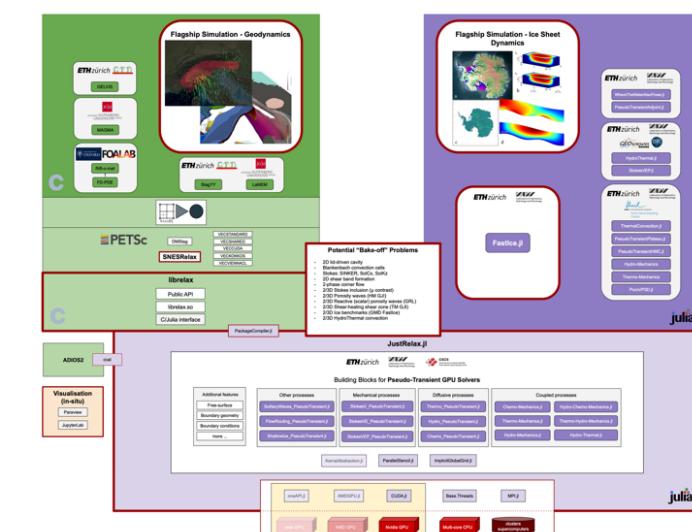
## GPU4GEO @ ETH / CSCS

- Computational earth science

...



Oceananigans.jl



Credit: Andrew R. Winters

# Julia HPC Community

## Julia Slack

- **#hpc**
- **#distributed**
- **#performance-helpdesk**



## JULIA FOR HPC

JuliaCon 2023 Minisymposium

## Monthly Julia HPC call

- <https://julialang.org/community/#events>

[arxiv.org/abs/2211.02740](https://arxiv.org/abs/2211.02740)

## At conferences

- JuliaCon23, SC23, PASC23,  
SIAM CSE 23, ...

### Bridging HPC Communities through the Julia Programming Language

Journal Title  
XX(X):1–14  
©The Author(s) 2022  
Reprints and permission:  
[sagepub.co.uk/journalsPermissions.nav](http://sagepub.co.uk/journalsPermissions.nav)  
DOI: 10.1177/ToBeAssigned  
[www.sagepub.com/](http://www.sagepub.com/)  
**SAGE**

Valentin Churavy<sup>1</sup>, William F Godoy<sup>2</sup>, Carsten Bauer<sup>3</sup>, Hendrik Ranocha<sup>4</sup>, Michael Schlottke-Lakemper<sup>5</sup>, Ludovic Räss<sup>6,7</sup>, Johannes Blaschke<sup>8</sup>, Mosè Giordano<sup>9</sup>, Erik Schnetter<sup>10,11,12</sup>, Samuel Omlin<sup>13</sup>, Jeffrey S. Vetter<sup>2</sup>, Alan Edelman<sup>1</sup>

#### Abstract

The Julia programming language has evolved into a modern alternative to fill existing gaps in the requirements of scientific computing and data science applications. Julia's single-language paradigm, and its proven track record at achieving high-performance without sacrificing user productivity, makes it a viable single-language alternative to the existing composition of high-performance computing (HPC) languages (Fortran, C, C++) and higher-level languages (Python, R, Matlab) suitable for data analysis and simulation alike. Julia's rapid growth in language capabilities, package ecosystem, and community make it a promising new universal language for HPC similar to C++ or Python – an achievable goal if the community is given the necessary resources. This paper presents the views of a multidisciplinary group of researchers in academia, government, and industry advocating for the use of Julia and its ecosystem in HPC centers. We examine the current practice and role of Julia as a common programming model to address major challenges in scientific reproducibility, data-driven artificial intelligence/machine learning (AI/ML), co-design, and in-situ workflows, scalability and performance portability in heterogeneous computing, network, data management, and community education. As a result, we consider necessary the diversification of current investments to fulfill the needs of the upcoming decade as more supercomputing centers prepare for the Exascale era.

#### Keywords

## Julia can be a great option for HPC!

- serial and parallel performance on-par with Fortran/C/C++
- portability and high-productivity
- new opportunities, e.g. interactive HPC



## New challenges

- workflow / interactivity at scale, Julia depot, system binaries, ...
- type instabilities

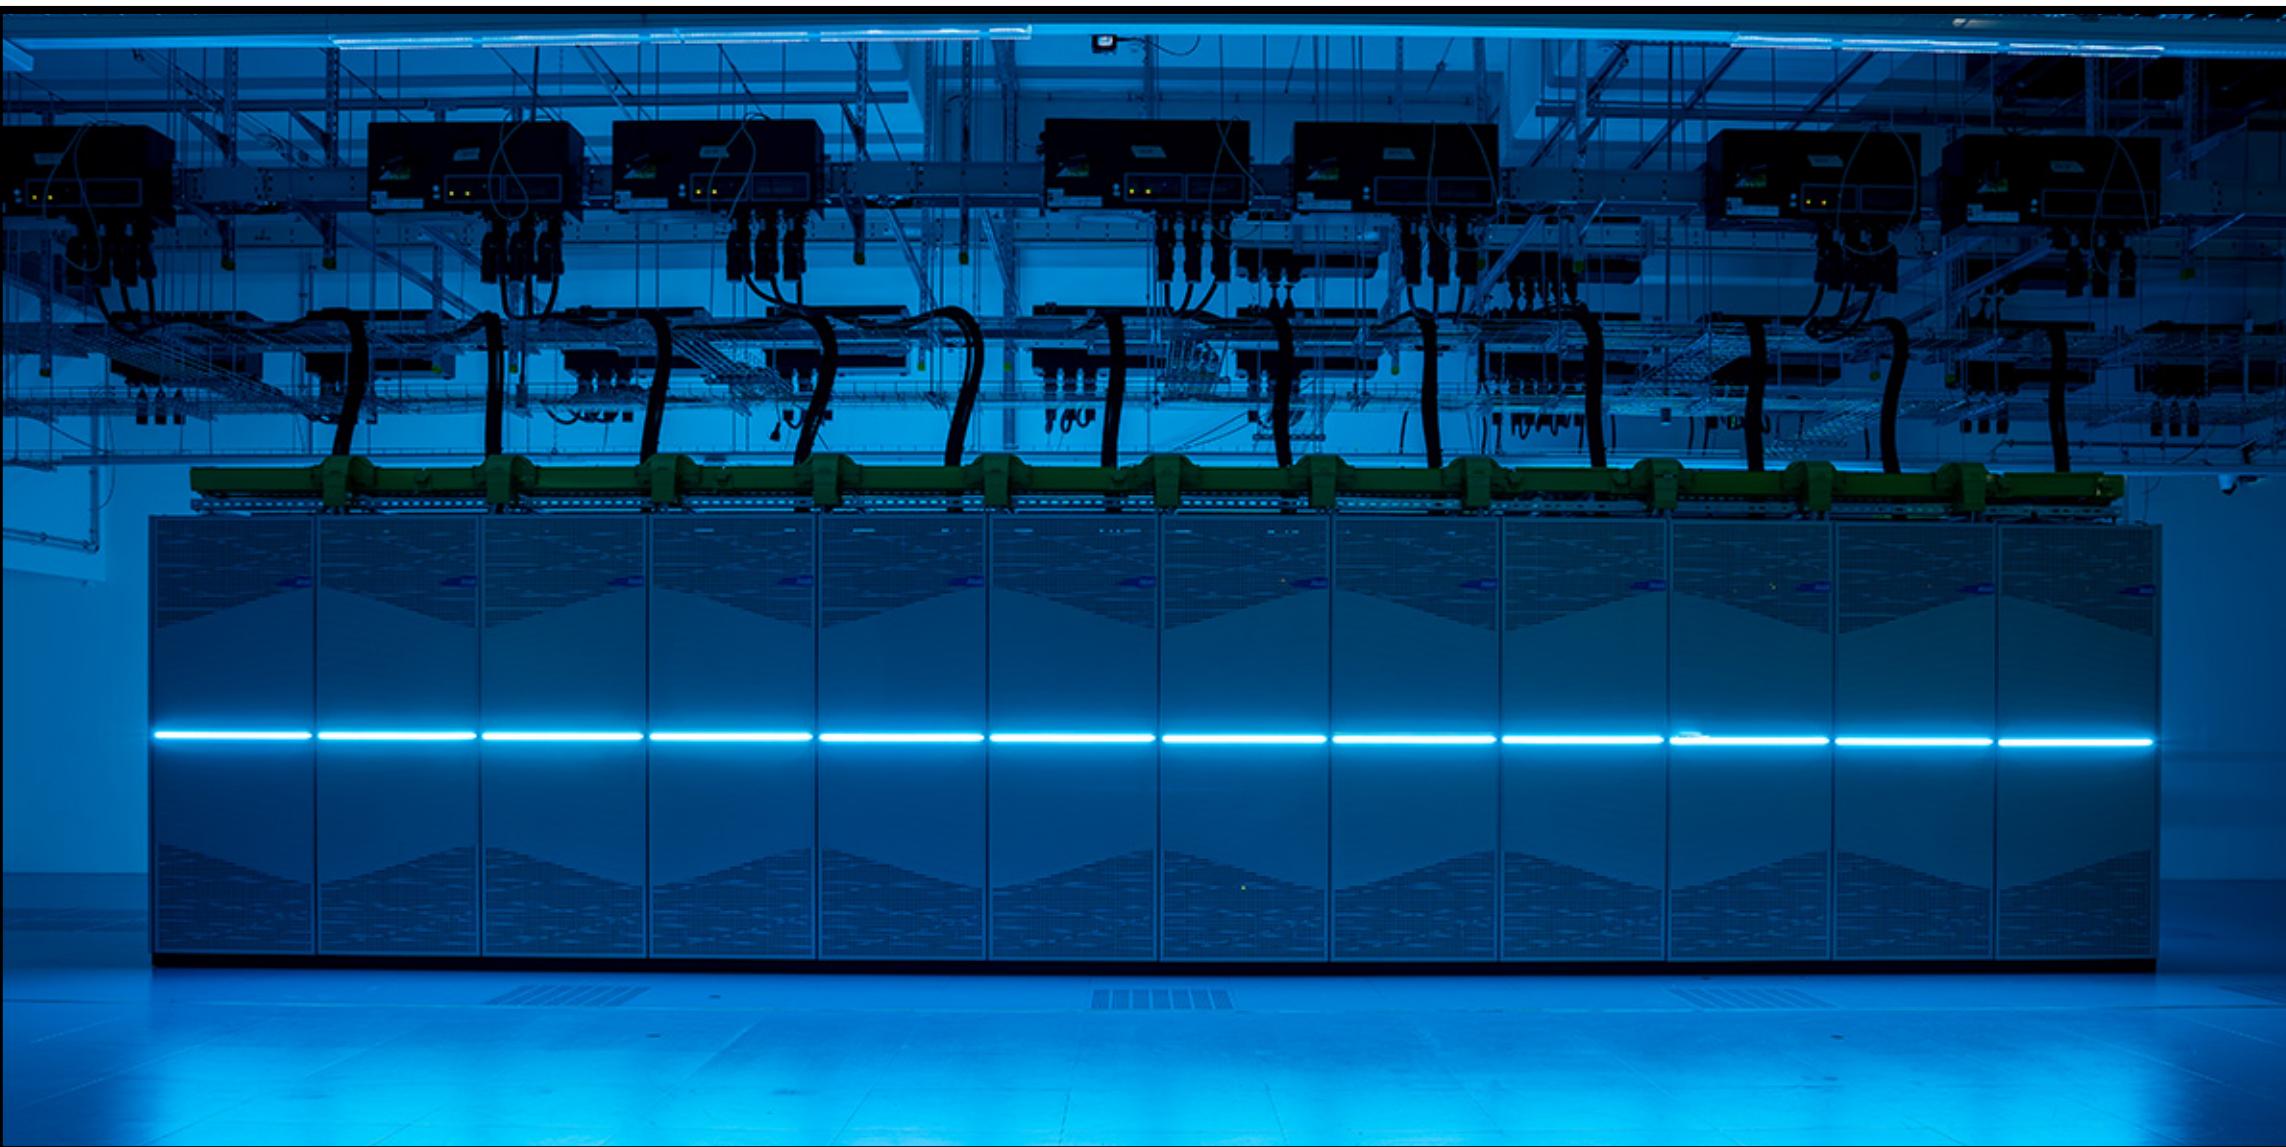
## “Early adoption” cost

- Julia for HPC is a niche
- lack of documentation, tooling, ...



# Let's get started!

- **PC2 JupyterHub**
  - <https://pc2.de/go/jupyterhub>



Noctua 2

