

# Onboarding

## Julia on HLRS Cluster/Laptops

September 10, 2024, HLRS

# HLRS Laptops

# Most of the course can be completed on the laptop.

- Equipped with NVIDIA GPU
- Jupyter + VS Code
  - `jupyter lab`
  - `code`
- Course materials
  - `$HOME/JuliaHLRS24`

# HLRS Training Cluster

# The training cluster has two kinds of compute nodes.

- CPU-only nodes “skl”
  - 2x Intel Skylake
  - 40 cores total
- GPU nodes “clx-ai”
  - 2x Intel Cascade Lake
  - 36 cores total
  - 8x **NVIDIA V100**

# HLRS uses the PBS Pro scheduler.

- Submit a job:
  - `qsub job_script.sh`
- See your queued/running jobs:
  - `qstat -rnw`

**VS Code → HLRS Cluster**

# Run VS Code on a cluster node via SSH

## Login node

- Works fine, just connect to (more later)
  - `accountname@training.hlrs.de`

## Compute node

- Possible but inconvenient
  - `SetEnv PBS_JOBID=2297805.hawk-pbs5`
  - `SSH ProxyJump`



# Side comment: “remote tunnels” instead of SSH

On the target node

- Download the `code` CLI and run
  - `code tunnel --verbose`

**Not at HLRS!**

Locally

- `Remote Tunnels: Connect to Tunnel`

# Julia on the Cluster

# Use a system module or standard Julia binaries.

- No need to compile Julia from source
- Binaries from
  - [juliaup](#) or [julialang.org](#)
- System module on [training.hlrs.de](#)
  - `module use julia`

# Put the Julia depot on the parallel file system.

- Why not \$HOME?
  - Quotas
  - Not always writable from compute jobs
- On HLRS clusters use **Workspaces**
  - `ws_allocate jlhpc 10`
- **JULIA\_DEPOT\_PATH** environment variable

# Use a Julia wrapper for the Julia VS Code extension.

- [Julia: Executable Path](#) should point to a wrapper script, like this one:

```
#!/bin/bash
[...]  
  
# Make julia available  
module load julia  
  
# Pass on all arguments to julia  
exec julia "${@}"
```

**Let's try it!**