Onboarding

Julia on HLRS Clusters/Laptops

Most of the course can be completed on the laptops.

- Equipped with NVIDIA GPU
- Jupyter + VS Code
  - `jupyter lab`
  - `code`
- Course materials
  - `$HOME/JuliaHLRS24`

The cluster training.hlrs.de has two types of nodes.

CPU nodes

- “skl”
- 2x Intel Skylake
- 40 cores total

GPU nodes

- “clx-ai”
- 2x Intel Cascade Lake
- 36 cores total
- 8x NVIDIA V100

Jobs are scheduled with PBS Pro.

- Submit a job:
  - `qsub job_script.sh`
- See your queued/running jobs:
  - `qstat -nw`

VS Code → HLRS Cluster

Run VS Code on a cluster node via SSH.

Login node

- Works fine, just connect to
  - `accountname@training.hlrs.de`

Compute node

- Possible but inconvenient
  - `SetEnv PBS_JOBID=...`
  - `SSH ProxyJump`

Can use “Remote tunnels” instead of SSH.

Not at HLRS!

On the target node

- Download the code CLI and run
  - `code tunnel --verbose`

Locally

- Remote Tunnels: Connect to Tunnel

Julia on the Cluster

Use a system module or standard Julia binaries.

- Modules on the HLRS training cluster
  - `module use julia`
  - `module use nvhpc # MPI+CUDA`
- Otherwise, use binaries provided by `juliaup`

Put the Julia depot on the parallel file system.

- Julia depot = where Julia stores stuff
  - packages
  - binary dependencies
  - ...
- Environment variable: `JULIA_DEPOT_PATH`
- Why not `$HOME`?
  - Quotas
  - Read-only from compute jobs (sometimes)

Need a Julia wrapper for the Julia VS Code extension.

- Julia: Executable Path should point to a wrapper script, like this one:

```
#!/bin/bash
[...]

# Load modules
module load julia
module load nvhpc

# Act like Julia (i.e. pass on all arguments)
exec julia "${@}"
```

Let’s do it!

→ exercises/Day1/1_cluster_onboarding/