

# Onboarding

## Julia on HLRS Clusters/Laptops

# Most of the course can be completed on the laptops.

- Equipped with NVIDIA GPU
- Jupyter + VS Code
  - jupyter lab
  - code
- Course materials
  - [\\$HOME/JuliaHLRS24](#)

# HLRS Training Cluster

# The cluster training.hlr.de has two types of nodes.

## CPU nodes

- “skl”
- 2x Intel Skylake
- 40 cores total

## GPU nodes

- “clx-ai”
- 2x Intel Cascade Lake
- 36 cores total
- 8x NVIDIA V100

# Jobs are scheduled with PBS Pro.

- Submit a job:
  - `qsub job_script.sh`
- Check on your queued/running jobs:
  - `qstat -nw`

**VS Code → HLRS Cluster**

# Run VS Code on a cluster node via SSH.

## Login node

- Works fine, just connect to
  - `accountname@training.hlrs.de`

## Compute node

- At HLRS, possible but inconvenient
  - `SetEnv PBS_JOBID=...`
  - `SSH ProxyJump`

**We will stay on Login nodes for the course.**

# Julia on the Cluster



# Use a system module or standard Julia binaries.

- Modules on the HLRS training cluster
  - `module use julia`
  - `module use nvidia/nvhpc` # MPI+CUDA
  - `module use compiler/nvidia` # MPI+CUDA
- If there is no (working) system module, use standard binaries provided by [juliaup](#).

# Put the Julia depot on the parallel file system.

- Julia depot = where Julia stores stuff
  - packages
  - binary dependencies
  - ...
- Environment variable: JULIA\_DEPOT\_PATH
- Why not \$HOME?
  - Quotas
  - Read-only from compute jobs (sometimes)

Already taken care of  
on HLRS training cluster!

# Need a Julia wrapper for the Julia VS Code extension.

- [Julia: Executable Path](#) should point to a wrapper script, like this one:

```
#!/bin/bash
[...]  
  
# Load modules  
module load julia  
module load nvidia/nvhpc  
Module load compiler/nvidia  
  
# Act like Julia (i.e. pass on all arguments)  
exec julia "${@}"
```

# Let's try it!

→ `exercises/Day1/1_cluster_onboarding/`

